

Task-Specific ESN Uncertainty Analysis

DROID-100 Task 5: Franka Panda Manipulation


Incremental learning curve proves generative model building.
Task-specific analysis gives 1.54x sigma ratio and AUC=0.785.

Author	Umar Bin Muzzafar
Organization	OMNIBIO LTD
Date	April 07, 2026
Dataset	DROID-100 Task 5 (Franka Panda, 35S + 19F)

The Task: What Is the Robot Doing?


DROID-100 Task 5 is a Franka Panda 7-DOF arm performing tabletop manipulation (pick-and-place with small objects). The dataset contains 54 episodes of this same task: 35 successful completions and 19 failures. Failed episodes are shorter on average (189 vs 364 frames) — the robot gets stuck or drops the object partway through.

Episode 4 (SUCCESS) — 213 frames



Successful episode (ep 4): full reach-grasp-move sequence

Episode 5 (FAILURE) — 87 frames



Failed episode (ep 5): shorter, robot struggles

Metric	Success	Failure
Episodes	35	19
Avg frames	364	189
Avg duration	24.3s	12.6s
State dims	7 joints	7 joints
FPS	15	15

Finding 1: Repeated Visits Decrease Uncertainty

We trained the ESN incrementally — first on 1 episode, then 2, then 3, up to 30 — and measured prediction error (sigma) on a held-out validation set after each step. If the ESN is learning a generative model, sigma should decrease as it sees more examples of the same task.




Figure 1: Learning curve. Left: validation sigma drops 94% as training data grows. Right: training RMSE stabilizes after ~15 episodes.

Sigma decreased by 94.3% from 9.29 (1 episode) to 0.53 (30 episodes). This confirms the ESN is building a generative model — each new episode of the same task makes the reservoir's predictions more accurate, analogous to increasing confidence in a Bayesian generative model.

Implications for BCM Learning

Currently the ESN reservoir weights are **fixed** — only the linear readout is trained. The learning curve shows this already works well. However, Rob suggested using **BCM (Bienenstock-Cooper-Munro) theory** to also adapt the reservoir weights online. BCM is a synaptic plasticity rule from neuroscience where:

- **Frequently active synapses get strengthened** (long-term potentiation) — states the robot visits often become better represented in the reservoir
- **Rarely active synapses get weakened** (long-term depression) — the reservoir forgets states it hasn't seen recently
- **The threshold between strengthening and weakening is adaptive** — it slides based on the neuron's recent activity history, preventing runaway excitation

Applying BCM to the ESN reservoir would let the weights adapt to the specific task's dynamics over time, potentially making the sigma signal even stronger. An alternative for Liquid State Machines (LSMs) is **Short-Term Plasticity (STP)**, which modulates synaptic strength on faster timescales. Both approaches let the reservoir specialize without catastrophic forgetting, because the adaptation is local and activity-dependent.

Finding 2: Task-Specific Sigma Separation

Training and evaluating the ESN within a single task (no cross-task mixing) produces a cleaner failure signal than the all-task analysis from Phase 1.

Metric	All-Task (Phase 1)	Task-Specific	Change
Sigma ratio (F/S)	1.35x	1.54x	+14%
AUC	0.701	0.785	+12%
F1 Score	0.49	0.70	+43%
Recall	0.74	0.84	+14%
Precision	0.37	0.59	+59%

Task-specific AUC improved from 0.701 to 0.785. At the optimal threshold ($\text{sigma_mean} \geq 0.2885$), we catch 84% of failures with 59% precision. This validates Rob's intuition: per-task training eliminates cross-task noise.




Figure 2: Task-specific ROC curve. AUC = 0.785.

Finding 3: Per-Joint Analysis (Task-Specific)

Within this single task, `joint_2` (upper arm/elbow) remains the most predictive joint at 2.64x failure/success ratio — even stronger than the 1.88x from the all-task analysis. This suggests that for this particular manipulation task, the upper arm positioning is where failures originate.




Figure 3: Per-joint prediction error for Task 5. Joint 2 has the highest failure/success ratio at 2.64x.

Note: Rob raised a valid point that per-joint findings may be task-dependent. This analysis confirms that — for this specific task, the upper arm joints dominate. Different tasks may stress different joints. The pipeline now supports per-task analysis to investigate this for any future task.

Finding 4: Early Warning at 27% Completion




Figure 4: Temporal sigma profile for Task 5. Failure sigma diverges at 27% completion.

Sigma diverges at 27% episode completion — the ESN detects trouble when the robot is only 27% through the task. 73% of the episode remains as potential intervention window. This is where **Curious Pi** and **Cautious Pi** policies would activate.

Next Steps: Curious Pi & Cautious Pi

When sigma exceeds the caution threshold, two policies can activate:

Cautious Pi ($\sigma > \text{caution threshold}$)

Slow down, reduce action magnitude, increase control gains. The robot becomes more careful when uncertainty is high. This is the conservative policy — minimize risk of failure.

Curious Pi ($\sigma > \text{curiosity threshold}$)

Perform exploratory actions designed to reduce sigma: nudge the object gently, lift slightly without committing, observe the state transition. The aim is to **update the generative model** — see how the state changes in response to small actions, making future predictions more accurate. This may involve BCM weight updates to the reservoir.

Both policies need the task-specific sigma threshold from this analysis ($\sigma_{\text{mean}} \geq 0.2885$) to trigger. The 27% early warning window gives the robot time to explore before committing.